

# *User Subroutines and the Runtime Stack*

## *CS 350: Computer Organization & Assembler Language Programming*

### **A. Why?**

- Subroutines are the most basic way to share executable code.
- Information for a procedure call is stored in an activation frame. At runtime, the activation frames form a stack (in C and C++) or heap (in Java).

### **B. Outcomes**

After this lecture, you should

- Understand how simple subroutines can be defined and used.
- What activation frames look like and how the runtime stack is used to save procedure call information.

### **C. Simple User-Written Subroutines**

- The LC-3 uses the JSR and JSRR commands to Jump to a SubRoutine.
- Both instructions set  $R7 \leftarrow PC$  before the jump so that the subroutine knows where to return to.
- JSR uses an 11-bit PC offset to find the address of the subroutine to go to:
  - $R7 \leftarrow PC; PC \leftarrow PC + \text{Sext}(\text{PCoffset11})$
- JSRR uses a base register to specify where to go to.
  - $\text{target} \leftarrow R[\text{Base}]; R7 \leftarrow PC; PC \leftarrow \text{target}$
  - (A subtle issue: When the base register is R7, we need the temporary variable to correctly swap PC and R7. If the base register is not R7, then the semantics above is equivalent to  $R7 \leftarrow PC; PC \leftarrow R[\text{Base}]$ .)



	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
JSRR	0	1	0	0	0	0	0	Base			0	0	0	0	0	0

- To return from a subroutine call, use JMP R7 (jump with R7 as the base register). The assembler also allows RET as a substitute mnemonic.

	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0
RET	1	1	0	0	0	0	0	1	1	1	0	0	0	0	0	0

### *D. Framework For a Simple Subroutine*

- Write comments that specify what registers should contain parameters, which ones will contain results, and which ones get modified but not restored.
- Begin by saving the registers you will modify and restore. These can include
- Registers you're using for intermediate calculations.
  - R0, if you're going to use any of the I/O traps (GETC and IN read a character into R0; OUT prints R0, and PUTS requires a pointer in R0).
  - R7, if you're going to call any other subroutines or a TRAP. (Executing JSR, JSRR, or TRAP will cause the value of R7 you need to return with to be overwritten.)
  - The easiest way to save registers is to store them into some variables set aside for that purpose. (This doesn't work for recursive subroutines.)
  - Alternatively, you can have semantics like "This routine may change R0" — then the onus of saving/restoring is on the user.
- Before you return from the subroutine, restore the registers you saved.
  - Then return using RET or JMP R7 (they're equivalent).
- Note that unless you save and restore R7, the JMP R7 will go to the instruction after the most recent JSR, JSRR, or TRAP in your subroutine. (If you're lucky, this will cause some sort of obvious problem — a bad calculation or an infinite loop.)

***E. Example: Multiply Subroutine***

- Our earlier multiplication program can be easily made into a subroutine. The first question to ask is the protocol for calling the routine. Since we're passing and returning numbers, it seems natural to use registers for that purpose.
- In `multiply_subr.asm`, the main program calls a `Mult` subroutine to set  $Z = X * Y$  by first setting  $R1 = X$ ,  $R2 = Y$ . Then it calls `Mult` to set  $R1 = R1 * R2$  and then stores  $R1$  into  $Z$ .

```
; multiply_subr.asm -- Multiplication as a subroutine
; ...
; Main program: Calculate X * Y, store result in product
;
        .ORIG      x3050
        LD         R1, X
        LD         R2, Y
        JSR        Mult                ; R1 = R1 * R2
        ST         R1, Z
        HALT

X        .FILL     12
Y        .FILL     8
Z        .BLKW     1                ; Holds X*Y at end
```

- It's important to comment the subroutine call interface.

```
; Multiply subroutine
; Set R1 = R1 * R2 (if R2 >= 0); set R1 = 0 if Y < 0.
;
; Property:
;   Y < 0 implies product = 0, k = Y
;   Y >= 0 implies product = X * (Y-k) and 0 <= k <= Y
;       (So when k = 0, product = X * Y)
;
; Register usage: R1 = X, R2 = Y; R3 = product, R4 = k
;
```

- The subroutine itself is defined using the framework already described, so it begins by saving the registers. Note all the labels for the routine begin with `MU`, just to keep them separate from the labels from other future routines we might add.

```
Mult    ST         R7, MU_R7        ; Save registers
        ST         R4, MU_R4
        ST         R3, MU_R3
```

- Then we initialize the product and counter and fall into the loop.

```

; Establish product and k properties
      AND      R3, R3, 0          ; product = 0
      ADD      R4, R2, 0          ; counter k = Y

; If Y < 0, product = 0; otherwise, maintain
; product = X * (Y - k) while k = Y, Y-1, ....
; Stop when k = 0
;
MU_Loop BRNZ    MU_Done          ; until k <= 0
      ADD      R3, R3, R1        ; product += X
      ADD      R4, R4, -1        ; k--
      BR       MU_Loop

```

- After the loop, we set the result register, restore the other registers, and return. Note we didn't actually change R7, so we didn't really need to save it, but it's easy to imagine a future change that might call a TRAP (to print an error message, for example.)

```

; After loop, R3 = X * Y (if Y >= 0) or 0 (if Y < 0)
; Set R1 to result, restore registers, and return
;
MU_Done ADD      R1, R3, 0        ; R1 = X * Y
      LD        R3, MU_R3        ; Restore registers
      LD        R4, MU_R4
      LD        R7, MU_R7
      JMP       R7              ; return to caller

; Save area for registers
;
MU_R1   .BLKW    1
MU_R3   .BLKW    1
MU_R4   .BLKW    1
MU_R7   .BLKW    1

```

### ***F. Example: Reading a String***

- The readstring program from the previous lecture can also be made into a subroutine fairly easily.

- It used R0, R1, R2, and R3, so we want to start by saving those registers in addition to R7. Instead of halting, we need to restore the registers and JMP R7 to return to the caller.
- We also need to establish a protocol for calling the routine — should the subroutine contain the buffer to read into? Or should the caller pass the address of the buffer? If it passes the address, how should we do this?
- Since the I/O traps GETC, IN, OUT, and PUTS use R0, it seems reasonable to have ReadLine to use R0; the user should set R0 to point to the buffer before calling ReadLine. Aside from saving and restoring the registers and renaming some labels, the rest of ReadLine is similar to the earlier version.
- The main program calls ReadLine twice. The first call shows how to call a close-by routine: It uses LEA to get the address of the routine and JSR to call it. The second call loads a pointer to ReadLine and uses JSRR to call it; this will work whether ReadLine is close by or far away from the subroutine call.

```
; read_subr.asm
;
; The main program exercises the ReadLine subroutine.
;
                .ORIG    x3000
                LEA      R0, msg1          ; read first message
                JSR      ReadLine

                LD       R0, addr_msg2     ; read second message
                LD       R7, addr_read     ;
                JSRR     R7
                HALT

; Buffers to read lines into
;
msg1            .BLKW    100
msg2            .BLKW    100

; Pointers &ReadLine and &msg2 to illustrate JSRR
;
addr_read      .FILL    ReadLine
addr_msg2      .FILL    msg2
```

- The ReadLine routine begins with usage comments and register saving.

```
; ReadLine: Read a return-terminated string into a
; buffer pointed to by R0. Uses the same pseudocode as in
; readstring.asm. Registers are restored before return.
;
; Register usage:
;   R0 = GETC/OUT char, char *R1 = bp (buffer position),
;   R2 = -(return char), R3 = temporary
;
```

```
ReadLine      ST      R0, RS_R0    ; Save R0
              ST      R1, RS_R1    ; Save R1
              ST      R2, RS_R2    ; Save R2
              ST      R3, RS_R3    ; Save R3
              ST      R7, RS_R7    ; Save R7
```

- Then we do some initialization and prompt for and read the first character

```
; Calculate R2 = -(ASCII char for return)
              LD      R2, RS_rc     ; R2 = return char
              NOT     R2, R2        ; R2 = -(return char) - 1
              ADD     R2, R2, 1     ; R2 = -(return char)

; Initialize bp = &buffer, prompt user for string,
; and read first char
;
              ADD     R1, R0, 0     ; bp = &buffer
              LEA     R0, RS_msg    ; get prompt message
              PUTS                      ; prompt for input
              GETC                      ; read char into R0
              ADD     R3, R0, R2    ; calculate R0 - return char
```

- The loop repeatedly echoes / stores / reads a character until we see '\n'. Then we end the string with a '\0' and print it.

```
; Repeat printing the char, storing it, and reading the next
; char until we read in a return char
;
RS_Loop       BRZ     RS_Done       ; until char read = return
              OUT                      ; print char read in
              STR     R0, R1, 0     ; *bp = char read in
              ADD     R1, R1, 1     ; bp++
              GETC                      ; read next char
              ADD     R3, R0, R2    ; calc char - return char
              BR      RS_Loop       ; continue loop
```

```

; When we see the end of the line, terminate the string
; in the buffer and print it out
;
RS_Done      OUT                ; echo the return char in R0
              AND      R3, R3, 0 ; get a null char ('\0')
              STR      R3, R1, 0 ; terminate buffer string
              LD       R0, RS_R0 ; point to buffer
              PUTS                ; print the string we read in
              LD       R0, RS_rc  ; get a newline
              OUT                ; end this line of output

```

- In addition to restoring the registers and returning, we have to declare some constants and a save area for the registers.

```

; Restore the registers and return
;
              LD       R7, RS_R7
              LD       R3, RS_R3
              LD       R2, RS_R2
              LD       R1, RS_R1
              LD       R0, RS_R0
              JMP      R7

RS_rc         .STRINGZ "\n"      ; ASCII newline char
RS_msg        .STRINGZ "Enter chars (then return): "

; Save area for registers
RS_R0         .BLKW  1
RS_R1         .BLKW  1
RS_R2         .BLKW  1
RS_R3         .BLKW  1
RS_R7         .BLKW  1

```

## ***G. Recursive Routines and the Runtime Stack***

- We know that saving call information with the code for the called routine doesn't work for recursive routines because the saved information would be overwritten by the next call.

- Let's study the following pseudocode for  $\text{factorial}(n)$ , annotated with positions A – E  
 $\text{fact}(n)$ :

(A) **if**  $n < 1$  **then**

$\text{returned value} = 1$  (B)

**else**

**call**  $\text{fact}(n-1)$ ;

(C) set  $\text{temp} = \text{callee's returned value}$

(D) our  $\text{returned value} = \text{temp} * n$

(E) **return**

- Each call of  $\text{factorial}$  gets its own **activation record** (or “frame”) that includes
  - local variable  $\text{temp}$  (set and used by callee)
  - local variables to save registers (set and used by callee)
  - return address (saved and used by caller)
  - returned value (value set by callee, used by caller)
  - parameter  $n$  (value set by caller)
- The **Runtime Stack** holds the frames; each call pushes a frame onto the stack; the top frame is for the active call; returning pops the frame from the stack.
  - The stack is usually implemented as a linked list, so each frame also includes a “dynamic link” pointer to the previous stack element.
  - In C/C++, on return, we deallocate the space for the local variables of the called routine. That's why in these languages a routine should not return a pointer to a local variables.
  - In Java, the space is not deallocated and the garbage collector will recover the space if necessary.
- Illustration: Call  $\text{fact}(3)$ , show state at (B) of base case call  $\text{fact}(0)$ 
  - (Note: Not bothering to show register save areas Also not showing link to top frame of stack.)



for fact(3)	for fact(2)	for fact(1)	for fact(0)
temp	temp	temp	temp
Dynamic link to frame for <code>main</code>	d.l. to frame for fact(3)	d.l. to frame for fact(2)	d.l. to frame for fact(1)
ret addr to (C)	ret addr to (C)	ret addr to (C)	ret addr to (C)
ret val	ret val	ret val	ret val = 1
n = 3	n = 2	n = 1	n = 0
Paused at C	Paused at C	Paused at C	Active at B

for fact(3)	for fact(2)	for fact(1)	for fact(0)
temp	temp	temp	
d.l. to frame for <code>main</code>	d.l. to frame for fact(3)	d.l. to frame for fact(2)	
ret addr to (C)	ret addr to (C)	ret addr to (C)	
ret val	ret val	ret val	ret val = 1
n = 3	n = 2	n = 1	n = 0
Paused at C	Paused at C	Active at C	Returned

for fact(3)	for fact(2)	for fact(1)
temp	temp	temp = 1
d.l. to frame for <code>main</code>	d.l. to frame for fact(3)	d.l. to frame for fact(2)
ret addr to (C)	ret addr to (C)	ret addr to (C)
ret val	ret val	ret val = 1
n = 3	n = 2	n = 1
Paused at C	Paused at C	Active at E

for fact(3)	for fact(2)
temp	temp = 1
d.l. to frame for main	d.l. to frame for fact(3)
ret addr to main pgm call	ret addr to (C)
ret val	ret val = 2
n = 3	n = 2
Paused at C	Active at E

for fact(3)
temp = 2
d.l. to frame for main
ret addr to (C)
ret val = 6
n = 3

Active at E

- In the example above, all the frames were for the same routine, so the frames were the same size.
- More generally, the size of a stack frame will depend on how many parameters and local variables the called routine has.

### ***General Method for Routine $P$ calling Routine $Q$***

- The caller  $P$  sets up the parameters of  $Q$  by calculating the argument values; then adds them to the activation we're building for  $Q$ .
- Then  $P$  does a JSR to the routine being called (saving  $R7 \leftarrow PC$  to establish the address to return to).
- The callee  $Q$  sets up the rest of the activation record;  $Q$  has to allocate space for

- The return value
- The return address (and set it to R7)
- The dynamic link (the pointer to the old top of the runtime stack)
- Any desired local variables, including for register saves
- We reset the top-of-runtime-stack pointer to the activation frame we just built.
- Then we jump to the code for the body of  $Q$  and start executing it.

### ***Returning from $Q$ to $P$ (in $C$ and $C++$ )***

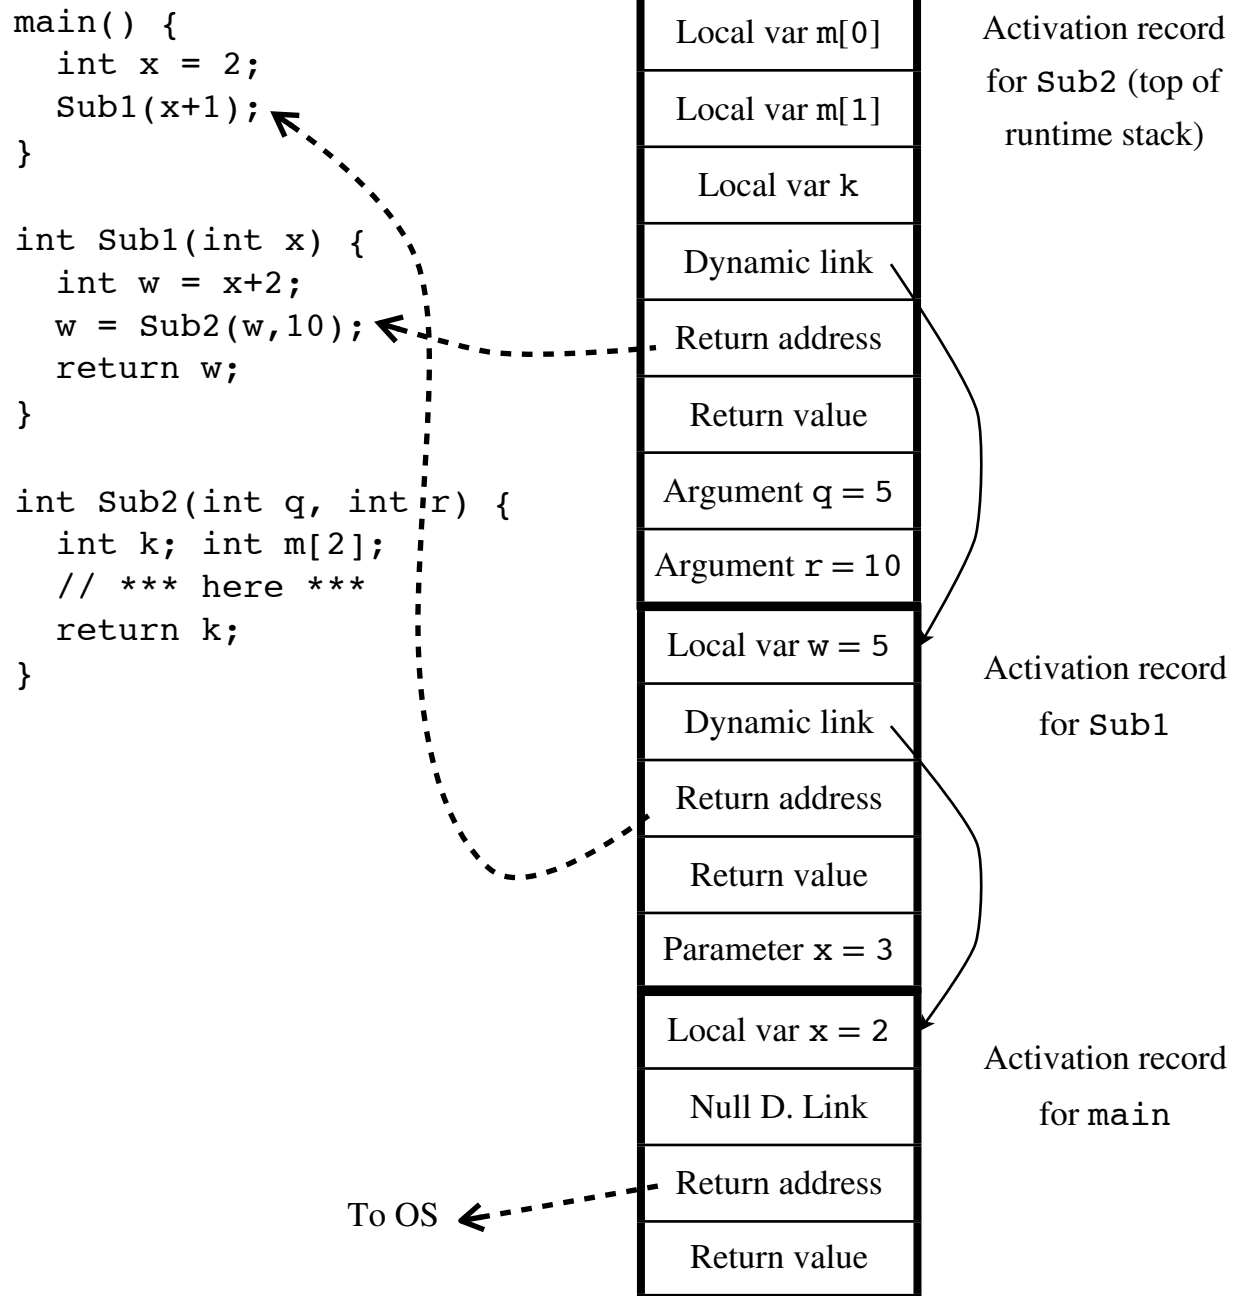
- For  $Q$  to return to  $P$ , we have to pop the activation record for its call.  $Q$  can:
  - Deallocate space for its local variables.
  - Get the dynamic link so we can reset the top of the runtime stack (and deallocate the space for the dynamic link).
  - Get the return address and restore R7 (and deallocate the space for the return address).
  - Copy the value being returned by  $Q$  (if there is one) into its designated space.
  - Use R7 to jump back to  $P$ .
- $P$  can get the returned value (if there is one) and deallocate the space for it.
- $P$  also has to deallocate the space for the arguments passed to  $Q$ .

### ***H. Sample of Nested Call***

- In the example below, the `main` program calls `Sub1` which calls `Sub2`:

```
main() {  
    int x = 2;  
    Sub1(x+1);  
}  
  
int Sub1(int x) {  
    int w = x+2;  
    w = Sub2(w,10);  
    return w;  
}  
  
int Sub2(int q, int r) {  
    int k; int m[2];  
    // *** Here ***  
    return k;  
}
```

- Note that since the `x` of the `main` program and the `x` of `Sub1` are in different routines, their values must be stored in different places.
  - (Note the runtime stack below is drawn vertically.)
- At the point marked “Here” in `Sub2`, the activation stack looks like this:



# ***LC-3 Subroutines and the Runtime Stack***

## ***CS 350: Computer Organization & Assembly Language Programming***

### ***A. Why?***

- Subroutines are the most basic way to share executable code.
- In operating systems, stacks can help track nested subroutine calls.

### ***A. Outcomes***

After this activity, you should be able to

- Understand how simple subroutines can be defined and used.
- Sketch the runtime stack at various points during the execution of a program.

### ***B. Problems***

1. Why do we have the called subroutine save/restore registers, not the calling routine?
2. In general, what behavior do you get if you call a subroutine that doesn't save and restore R7 before calling a TRAP or subroutine?
3. Given the save/restore register technique we used in the sample programs, what happens if you try to write a recursive subroutine?
4. Sketch the code for a simple subroutine that assumes R0 points to a string, finds the length of the string, and returns the length in R1.
5. Suppose in the Sub1/Sub2 example, Sub2 returns 8. Sketch the runtime stack after Sub2 returns to Sub1 and Sub1 assigns `w = Sub2 ( ... )`.
6. Continuing from Question 1, sketch the runtime stack just after Sub1 returns to the `main` program.
7. If a subroutine wants to pass an array instead of an integer, then in C-like languages, we just copy the address of the initial element. What would the pros and cons be of copying the entire array onto the stack instead?

8. If routine  $P$  calls routine  $Q$ , which of the following would you find or not find in the activation record for the call?
- (a) The return address to the code for  $P$ .
  - (b) The return address from  $P$  back to its caller.
  - (c) A link to the activation record for  $P$ .
  - (d) A link to the activation record for  $Q$ .
  - (e) Space for the local variables of  $P$ .
  - (f) Space for the local variables of  $Q$ .
  - (g) Space for the arguments being passed to  $Q$ .
  - (h) Space for the arguments  $P$  received from its caller.
  - (i) Space for the value that  $Q$  will return to  $P$ .
  - (j) Space for the value that  $P$  will return to its caller.